

Hybrid Fine Grained ILU

Samah Karim
Peter Sentz
Shelby Lockhart

University of Illinois at Urbana-Champaign

Outline

- 1 Iterative Methods
- 2 Incomplete Factorization
- 3 New Parallel ILU Algorithm
- 4 Hybrid Fine Grained ILU
- 5 Current Project State

Outline

- 1 Iterative Methods
- 2 Incomplete Factorization
- 3 New Parallel ILU Algorithm
- 4 Hybrid Fine Grained ILU
- 5 Current Project State

Iterative Methods

- Direct methods for solving $Ax = b$ can be unnecessarily expensive
- Iterative methods are "infinite" solution techniques that requires computation of $x \rightarrow Ax$
- Can sometimes be computationally inefficient if too many iterations are required
- Their convergence can be accelerated by preconditioning

Preconditioning

$$Ax = b$$

Apply preconditioner

$$C^{-1}Ax = C^{-1}b$$

Preconditioner

A preconditioner is a matrix C such that:

- $\text{cond}(C^{-1}A) < \text{cond}(A)$
- Solution of linear systems $Cz = y$ requires little computational effort
- That is, computation of $y \rightarrow C^{-1}y$ is "easy"

Challenges

- If we choose $C = I$
 - no cost associated with solving linear system $Cz = y$
 - condition number is the same
- If we choose $C = A$
 - resulting $C^{-1}A = I$ perfectly conditioned
 - cost required is equal to that of solving the initial linear system with A itself
- Cost of computing C should be taken into account

No Perfect Answer!

- Optimal C minimizes total computational effort
- Computational effort depends on:
 - Size of the problem
 - Sparsity pattern of A
 - Spectrum of A
- No general answer about how to choose an optimal C
 - One method is Incomplete Factorization

Incomplete Factorization

- Given general sparse matrix A , an incomplete LU (ILU) factorization is a pair of matrices:
 - Sparse lower triangular matrix L
 - Sparse upper triangular matrix Usuch that $LU \approx A$
- Take preconditioner $C = LU$. Computing $C^{-1}y = U^{-1}L^{-1}y$ is simple:
 - Solve $Lw = y$ by forward substitution
 - Solve $Uv = w$ by backward substitution

This paper

- New fine-grained parallel algorithm for ILU factorization
- Computes all nonzero elements in parallel
- Iteratively improves the accuracy of the factorization
- Provides large amounts of parallelism using shared memory

Outline

- 1 Iterative Methods
- 2 Incomplete Factorization**
- 3 New Parallel ILU Algorithm
- 4 Hybrid Fine Grained ILU
- 5 Current Project State

Computing Incomplete Factorization

- Given general sparse matrix A , an incomplete LU (ILU) factorization computes:
 - Sparse lower triangular matrix L
 - Sparse upper triangular matrix Usuch that $LU \approx A$
- Use Gaussian elimination process but only allow nonzeros in specified locations of L and U
- Define sparsity pattern S such that $(i, j) \in S$ if:
 - l_{ij} is permitted to be nonzero when $i \geq j$
 - u_{ij} is permitted to be nonzero when $i \leq j$


Incomplete Factorization

Algorithm 1 Conventional ILU Factorization

```
1: for  $i = 2$  to  $n$  do
2:   for  $k = 1$  to  $i - 1$  and  $(i, k) \in S$  do
3:      $a_{ij} = a_{ik} / a_{kk}$ 
4:     for  $j = k + 1$  to  $n$  and  $(i, j) \in S$  do
5:        $a_{ij} = a_{ij} - a_{ik} a_{kj}$ 
6:     end for
7:   end for
8: end for
```

Example: Zero fill-in ILU (ILU(0))

Take S exactly as the sparsity pattern of A



Outline

- 1 Iterative Methods
- 2 Incomplete Factorization
- 3 New Parallel ILU Algorithm**
- 4 Hybrid Fine Grained ILU
- 5 Current Project State

New Parallel ILU Algorithm

- Parallelism is very fine-grained
- Individual entries of the factorization can be computed in parallel, using an iterative algorithm

Reformulation of ILU

- New algorithm based on the sometimes overlooked property that:

$$(LU)_{ij} = a_{ij}, \quad (i, j) \in S \quad (1)$$

where $(LU)_{ij}$ is the (i, j) entry of the ILU factorization of the matrix with entries a_{ij}

- i.e. factorization is exact on the sparsity pattern S

Reformulation of ILU

- New fine-grained algorithm interprets an ILU factorization as a problem of computing unknowns
 - l_{ij} , $i > j$, $(i, j) \in S$
 - u_{ij} , $i \leq j$, $(i, j) \in S$using property (1) as constraint.
- \Rightarrow total number of unknowns is $n_S = |S|$, the number of elements in the sparsity pattern S

Reformulation of ILU

- To find these n_S unknowns, rewrite constraints (1) as:

$$\sum_{k=1}^{\min(i,j)} l_{ik} u_{kj} = a_{ij}, \quad (i, j) \in S \quad (2)$$

- Each constraint can be associated with an element in $S \Rightarrow$ there are n_S constraints
- Thus problem of solving for n_S unknowns using n_S nonlinear equations

Solution of constraint equations

- Write explicit formula for each unknown in terms of the other unknowns:

$$l_{ij} = \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), \quad \text{if } i > j \quad (3)$$

$$u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, \quad \text{if } i \leq j \quad (4)$$

- Equations are in the form $x = G(x)$ where x is a vector containing the unknowns l_{ij} and u_{ij}

Solution of constraint equations

- Solve equations via fixed-point iteration:

$$x^{(p+1)} = G(x^{(p)}), \quad p = 0, 1, \dots \quad (5)$$

with an initial guess $x^{(0)}$

- Each component of new iterate $x^{(p+1)}$ can be computed in parallel

Initial Guess

- To begin the fixed point iterations, we need an initial guess for the unknowns
- Simple initial guess:
 - $(L^{(0)})_{ij} = (A)_{ij} \quad , i > j \quad , (i, j) \in S$
 - $(U^{(0)})_{ij} = (A)_{ij} \quad , i \leq j \quad , (i, j) \in S$

New Parallel ILU

Algorithm 2 Fine-Grained Parallel Incomplete Factorization

```

1: Set unknowns  $l_{ij}$  and  $u_{ij}$  to initial values
2: for  $sweep = 1, 2, \dots$  until convergence do
3:   for  $(i, j) \in S$ , in parallel do
4:     if  $i > j$  then
5:        $l_{ij} = (a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj}) / u_{jj}$ 
6:     else
7:        $u_{ij} = a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}$ 
8:     end if
9:   end for
10: end for

```

New Parallel ILU

- Each fixed-point iteration updating all unknowns is called a sweep
- Algorithm parallelized across elements of S
- Given p compute threads, partition set S into p parts, one for each thread
- Threads run in parallel updating the components of the vector of unknowns x , asynchronously

Outline

- 1 Iterative Methods
- 2 Incomplete Factorization
- 3 New Parallel ILU Algorithm
- 4 Hybrid Fine Grained ILU**
- 5 Current Project State

Distributed Memory Implementation

- To take advantage of available computing resources, algorithm should be adapted to distributed memory systems
- May want to increase problem size, where a matrix may no longer fit in a single node's memory

Block Jacobi Preconditioner


- A simple preconditioner is the diagonal of A , i.e. use $D = \text{diag}(a_{11}, a_{22}, \dots, a_{nn})$
- Can generalize this if we write A in *block* format:

$$A = \begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1K} \\ A_{21} & A_{22} & \cdots & A_{2K} \\ \vdots & \vdots & \ddots & \vdots \\ A_{K1} & A_{K2} & \cdots & A_{KK} \end{bmatrix}$$

- Use the *block diagonal* of A as a preconditioner:

$$D = \begin{bmatrix} A_{11} & 0 & \cdots & 0 \\ 0 & A_{22} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & A_{KK} \end{bmatrix}$$

Block Jacobi - Hybrid Implementation



Hybrid Implementation

- Distribute blocks across processors/nodes
- Compute ILU factorization for each block using OpenMP on each node
- Gather blocks on single processor in order to use preconditioner in iterative method
- Alternatively, could perform iterative method in parallel - no need to gather blocks

Outline

- 1 Iterative Methods
- 2 Incomplete Factorization
- 3 New Parallel ILU Algorithm
- 4 Hybrid Fine Grained ILU
- 5 Current Project State**

Work so far

- Have finished shared memory ILU implementation using OpenMP
- Have finished PCG (an iterative method) code utilizing forward/backward substitution and matrix multiply routines
- Currently finishing Hybrid Block Jacobi-ILU implementation using MPI across nodes

References

- [1] Edmond Chow and Aftab Patel. Fine-grained parallel incomplete lu factorization. *SIAM Journal on Scientific Computing*, 37(2):C169–C193, 2015.
- [2] Yousef Saad. *Iterative methods for sparse linear systems*. SIAM, 2003.