

Programming Techniques for Supercomputers:

Parallel Computers

Introduction

Shared-memory computers

Distributed-memory computers / Hybrid systems

Networks - Introduction

Prof. Dr. G. Wellein^(a,b) , Dr. G. Hager^(a) , M. Kreutzer^(a)

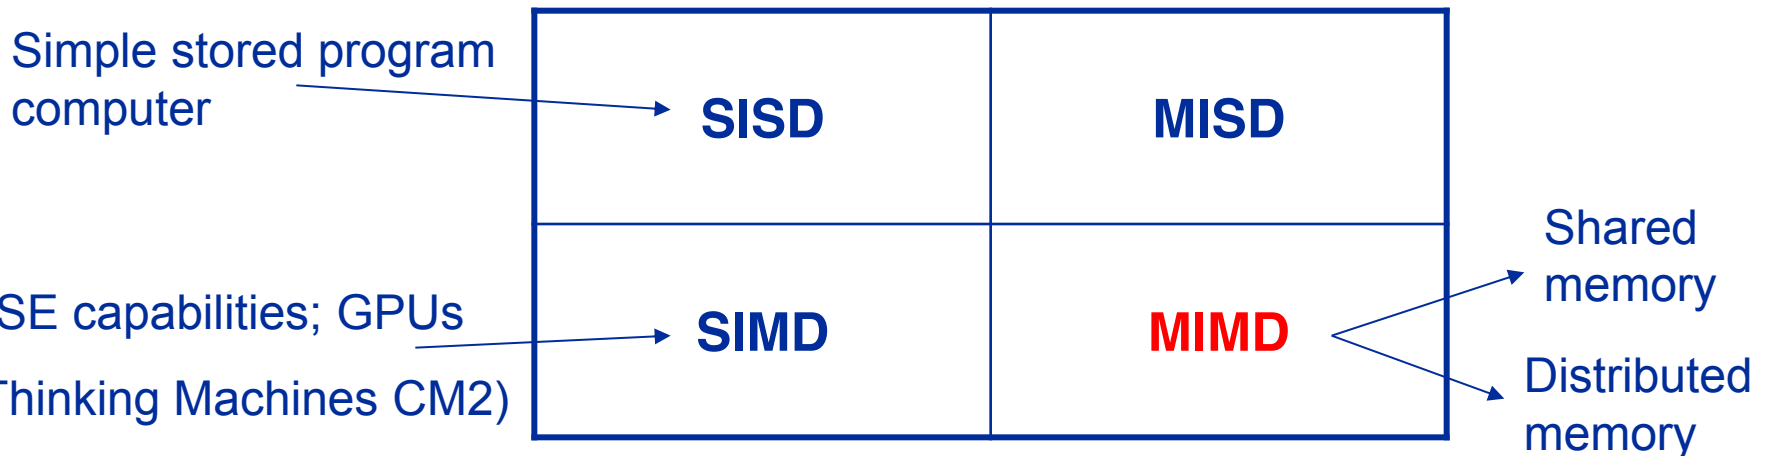
^(a)HPC Services – Regionales Rechenzentrum Erlangen

^(b)Department für Informatik

University Erlangen-Nürnberg, Sommersemester 2013




- „Parallel Computing“ : A number of compute elements solve a problem in a cooperative way
- **Parallel Computer:** A number of compute elements connected such way to do parallel computing for a large set of applications
- **Classification according to Flynn:** **M**ultiple **I**nstruction **M**ultiple **D**ata (**MIMD**)





- **Shared memory computers provide**

- a single shared address space (memory) for all processors
- All processors share the same view of the address space!

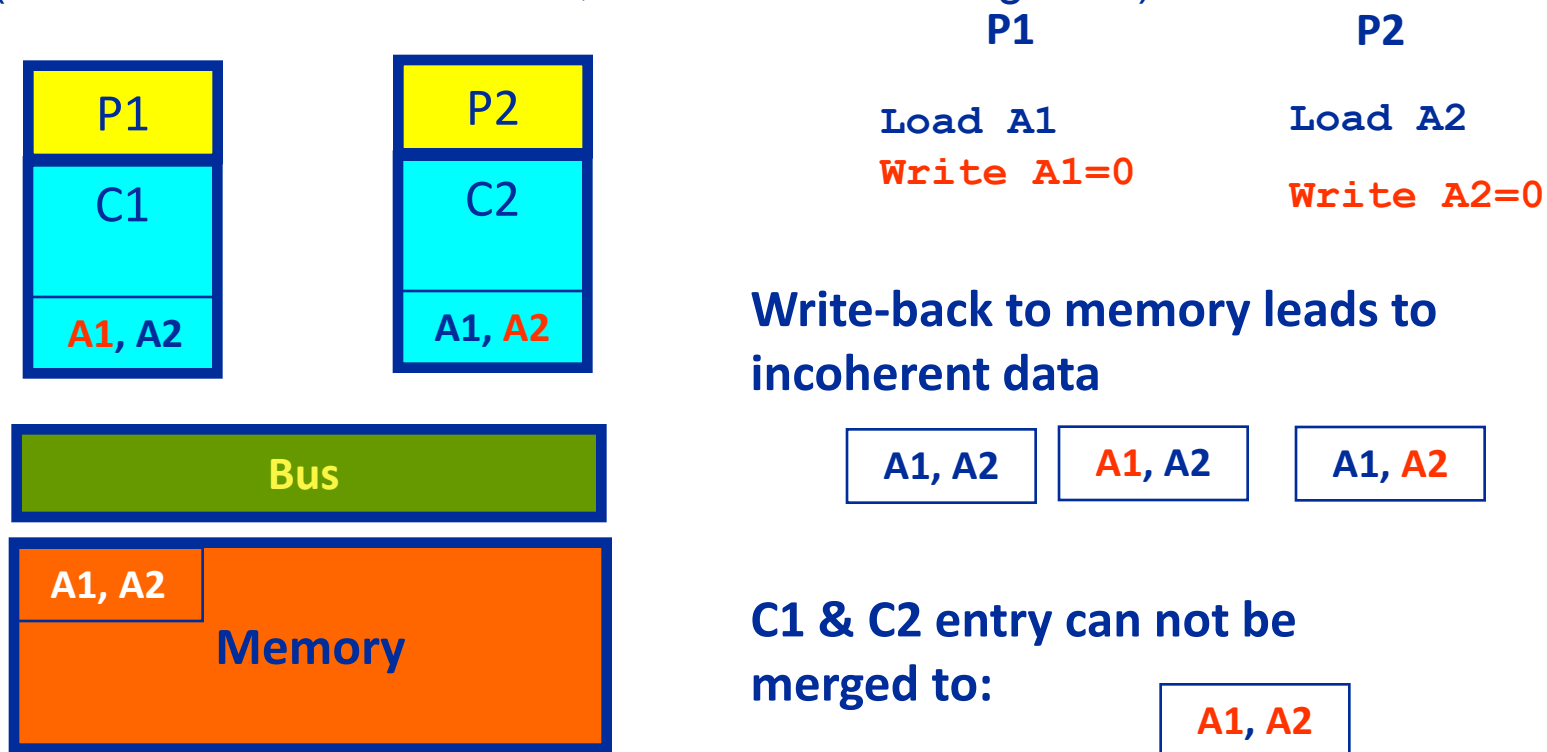


- **Two basic categories of shared memory systems**

- **Uniform Memory Access (UMA):**
Memory is equally accessible to all processors with the same performance (Bandwidth & Latency)
- **cache-coherent Non Uniform Memory Access (ccNUMA):**
Memory is physically distributed but appears as a single address space: Performance (Bandwidth & Latency) is different for local and remote memory access
- Copies of the same cache line may reside in different caches → Cache coherence protocols guarantees consistency all time (for UMA & ccNUMA)
- Cache coherence protocols do not alleviate parallel programming for shared-memory architectures!

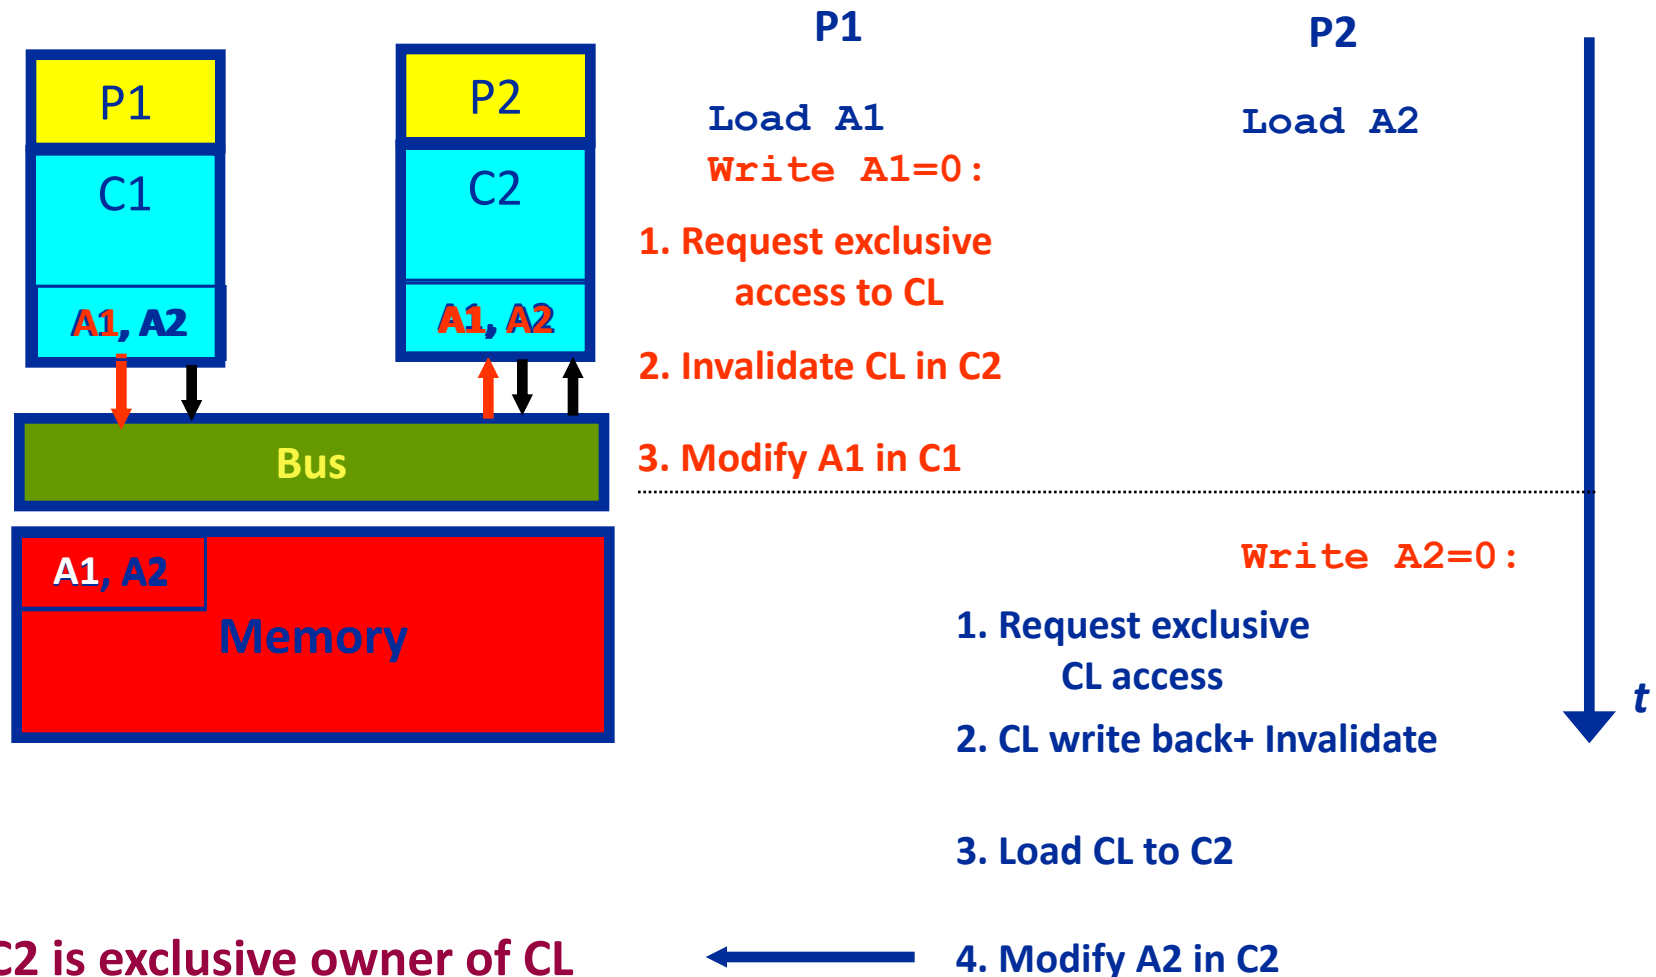


- **Data in cache is only a copy of data in memory**
 - Multiple copies of same data on multiprocessor systems
 - Cache coherence protocol/hardware ensure consistent data view
 - Without cache coherence, shared cache lines can become clobbered:
(Cache line size = 2 WORD; A1+A2 are in a single CL)





- Cache coherence protocol must keep track of cache line status





- **Cache coherence can cause substantial overhead**
 - may reduce available bandwidth
- **Different implementations**
 - **Snoop**: On modifying a CL, a CPU must broadcast its address to the whole system
 - **Directory, “snoop filter”**: Chipset (“network”) keeps track of which CLs are where and filters coherence traffic
- **Directory-based ccNUMA can reduce pain of additional coherence traffic**
- **But always take care:**

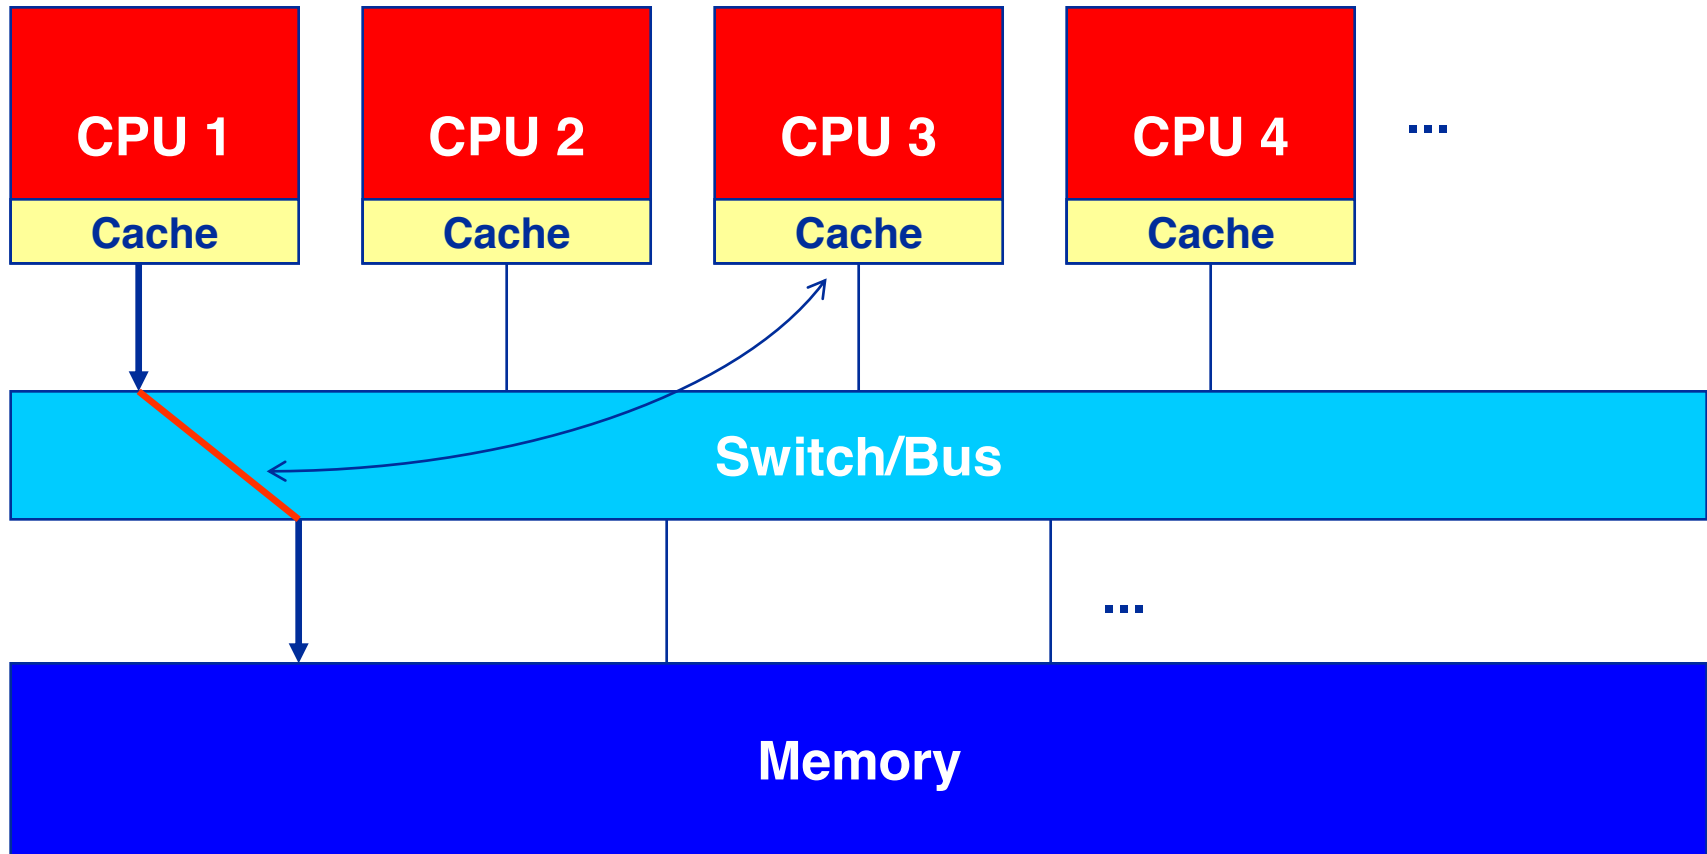
Multiple processors should never write frequently to the same cache line (“false sharing”)!



- Widespread cache coherence protocol: **MESI** protocol
- A cache line can have four different states:
 - **M**odified: Cache line has been modified in this cache, and it resides in no other cache. Cache line needs to be evicted to ensure memory consistency
 - **E**xclusive: Cache line has been read from main memory but not (yet) modified. There are no (valid) copies in other caches
 - **S**hared: Cache line has been read from memory but not modified. There may be valid copies in other caches
 - **I**nvalid: This cache line does not reflect any sensible data. Usually this happens if the cache line was in **S** state and another processor request exclusive ownership





- **UMA Architecture: switch/bus arbitrates memory access**
 - Special protocol ensures cross-CPU cache data consistency
 - Flat memory – also known as „Symmetric Multi-Processor“ (SMP)






- **Worst case: bus system provides single bandwidth to multiple processors**
 - Only one CPU at a time can use the bus and access memory at any one time – No need to provide for faster memory
 - Collisions occur frequently, causing one or more CPUs to wait for “bus ready” (contention)
 - Even worse: Shared memory bus in current multi-core chips further reduce available bandwidth/peak FLOP balance!
- Price/performance ratio was good for certain applications for a long time
- Similar concept in a single multi-core chip





- **Best case: memory crossbar switch** provides separate data path to memory for each CPU
 - Can saturate full memory bandwidth of every CPU concurrently
 - Bus contention occurs only if same memory module is accessed by more than one CPU
 - Memory interleaving is a must
 - Example:
 - NEC SX-9 node: 16 CPUs, 256 GB/sec each, 4096 GB/sec per node bandwidth
- **Reality: Compromises are made**
 - E.g. 2 of 8 cores can use full bandwidth concurrently
 - May be advantageous not to use all cores on a node
 - Example: Intel Clovertown based on 2 quad-core chips 2 sockets (8 cores), 10.6 GB/sec each, 21.2 GB/sec node bandwidth; i.e. 2.66 GB/s per core





- **Examples:**

- Your dual-/quad-/hexa-core laptop/desktop computer
- IBM BlueGene series
- NEC vector systems

- **Advantages**

- Cache Coherence (see below) is "easy" to implement
- Easy to optimize memory access
- Incremental parallelization
- Large memory configuration in a single address space

- **Disadvantages**

- Memory bandwidth and price (!) often limit scalability (2 – 8 cores per UMA node)



- **ccNUMA:**

- Single address space although physically distributed memory through proprietary hardware concepts (e.g. NUMALink in SGI systems; QPI for Intel; HT for AMD)


- **Advantages:**

- Aggregate memory bandwidth is scalable
- Systems with more 1024 cores are available (SGI)

- **Disadvantages:**

- Cache Coherence hard to implement / expensive
- **Performance depends on access to local or remote memory**

- **Examples: All modern multi-socket compute nodes – SGI Altix/UV**





■ System Architecture cont'd

- Layout for **4-socket AMD Opteron** system $\rightarrow 4 \times (2 \times 8 \text{ cores}) = 64 \text{ cores}$ with **8 NUMA domains**
- ... and you know this can get really large!

■ Who keeps track of the contents of all those caches?

- **Cache coherence mechanisms**

- How do we make sure that memory pages are as close as possible to the CPUs they are needed on?
 - **ccNUMA optimization**

