

CS 484 SP17
Midterm 2, April 5th, 2017

Total time: 75 mins
No outside materials or devices allowed

Name: _____

NetID: _____

Question	Total Points	Points Obtained
Question 1	10	
Question 2	15	
Question 3	15	
Question 4	10	
Question 5	10	
Question 6	15	
Total Points	75	

1. True or False / Multiple Choice: Multiple answers may be correct for some questions.
(2pts each)

a) An algorithm with a larger isoefficiency function (e.g. cubic) is more efficient as the problem size grows than an algorithm with a smaller (say quadratic) isoefficiency function.	True / False
b) Charm++ divides the work between the system and the programmer such that the programmer can decide how to balance the load, while the runtime system decomposes the computation into parallel pieces.	True / False
c) By specifying PUP method for your chare class, you are providing the runtime system information that allows it to:	A) Decide the priority of the chare B) Decide which processor to migrate the chare C) Decide how to copy the object's data into a buffer for migration D) Decide which method to invoke after load balancing is completed to resume execution
d) The efficiency of a parallel program is given by which equation? Assume E = efficiency, S = sequential time, P = parallel time, and N = number of processors.	A) $E = S / (N * P)$ B) $E = P / (N * S)$ C) $E = S / (N + P)$ D) $E = P / (N + S)$
e) Which of the following sorting algorithms can be parallelized?	A) Radix Sort B) Histogram Sort C) Quicksort D) All of the above
f) There are no blocking entry methods in Charm++, i.e., all entry methods are asynchronous and can only have "void" as their return type.	True / False
g) Chare method invocations are asynchronous.	True / False
h) Which MPI call would you use to combine values from all processes and distributes the result back to all processes?	A) MPI_Allreduce B) MPI_Broadcast C) MPI_Scatter D) MPI_Alltoall

2. Short Answer:

(5+5+5 points)

Consider the following charm++ program where E is the entry method of a chare. A_Proxy and B_Proxy are proxies to two distinct 1-dimensional chare arrays, The programmer's intent is: send x to the foo method of the 23rd element of the chare array referenced by A_Proxy, and calculate y. E then gets y and sends it to all elements of the chare array referenced by B_Proxy.

```
void E(...) {  
    ...  
    A_Proxy[23].foo(x, &y);  
    B_Proxy[ALL].bar(y);  
    ...  
}
```

A. Explain all the errors the programmer has made in the code provided. (Assume all the necessary variables, proxies and entry methods are declared and initialized)

Errors :

Methods invocations in charm++ are asynchronous, there is no guarantee that the call to A_proxy will be completed before the calls to B_proxy. The index of A_proxy should be 22 since we are counting from 0. ALL is not a valid keyword for invoking a method on all chares in a chare array.

B. Suggest a solution to fix the issues. You are allowed to change the signatures of *foo*. Full points only if you don't add a new entry method, but correct solutions with additional entry method will still get substantial credit. Describe in pseudocode what *foo* should do.

Correction :

One possible solution

Pass B_proxy as a parameter to foo, compute y and use the result as parameter to the bar method of B_proxy

```
A_proxy[22].foo(x,&y,B_proxy b)
{
    y = .....;
    b.bar(y);
}
```

C) Assume a parallel algorithm with problem size W requires N^3 operations. The amount of computation on each processor is equal to $\frac{N^3}{P}$. The amount of communication is given by the equation $4 * N$. Compute the isoefficiency of the algorithm.

$$W = N^3$$

$$T_c = \frac{N^3}{P}$$

$$T_p = 4 * N$$

$$T_t = T_c + T_p$$

$$\begin{aligned} T_o &= p * T_t - N^3 \\ &= p * 4 * N + N^3 - N^3 \\ &= 4 * p * N \end{aligned}$$

$$W = K * T_o$$

$$W = K * 4 * p * N$$

$$N^3 = K * 4 * p * N$$

$$N^2 = K * 4 * p$$

$$= 2\sqrt{Kp}$$

$$\begin{aligned} W &= K * 4 * p * 2 * \sqrt{Kp} \\ &= 8 * (Kp)^{\frac{3}{2}} \end{aligned}$$

3. MPI + OpenMP

(15 pts)

A. Using MPI and OpenMP in conjunction, finish the program below for computing the sum of $f(n)$ for each number n in a distributed array *globalArray*. For simplicity, assume $f()$ and $calculateithValue()$ are defined elsewhere, and that the total length of the array is exactly divisible by the number of processes.

```
int main() {
    //Initialize MPI stuff
    int world_size, rank, provided;
    MPI_Init_Thread(NULL, NULL, MPI_THREAD_FUNNELED, &provided);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    //Initialize arrays and variables for computation
    const static int array_length = MAX_LENGTH;
    int localLength = array_length/world_size;
    int* localArray = (int*)malloc(sizeof(int)*localLength);
    int *globalArray;
    int globalSum = 0, local_sum = 0;

    //Initialize the array for the first process
    if (rank == 0)
    {
        globalArray = (int*)malloc(array_length*sizeof(int));
        for(int i=0;i<array_length;i++)
            globalArray[i] = calculateithValue(i);
    }

    //Write parallel code that accomplishes the same thing as the
    // following sequential code. Assume the computations of f() are
    independent :
    int sequentialSum = 0;
    for(int i = 0; i < MAX_LENGTH; ++i)
        sequentialSum += f(globalArray[i]);
}
```

//Write MPI code here to distribute process 0's globalArray to all other processes' localArrays, then compute partial sums on each process using OpenMP. Finally, add all processes' local_sum's together using MPI to compute globalSum. The final globalSum should be available on process 0.

```
MPI_Scatter(&global_array[0], localLength, MPI_INTEGER, &localArray[0], localLength, MPI_INTEGER, 0, MPI_COMM_WORLD);
```

```
int local_sum = 0;
```

```
#pragma omp parallel num_threads(nth)
{
    #pragma omp for
    for(int i=0; i<localLength; i++)
    {
        #pragma omp critical
        local_sum += localArray[i];
    }
}
```

```
int global_sum = 0;
```

```
MPI_Reduce(&globalSum, &local_sum, 1, MPI_INTEGER, MPI_SUM, 0, MPI_COMM_WORLD);
```

```
if (rank == 0)
    assert(globalSum == sequentialSum);
```

```
//Clean up
free(localArray);
MPI_Finalize();
return 0;
}
```

4. Shared-Memory Atomics

(10 pts)

Consider the following sequential program. Assume the buffer is bounded, and the size of the buffer is larger than the number of threads. Use c++ atomics for synchronization.

```
int *s = (int *)malloc(MAX_SIZE*sizeof(int));
assert ( s != NULL);

// Keep this loop sequential
for(int i=0;i<MAX_SIZE;i++) s[i] = rand()%1000;

// Re-write this sequential section using threads (pthreads/c++
threads)

int histogram[RANGE];
for(int i=0;i<MAX_SIZE;i++) histogram[s[i]]++;

//Write down your parallel solution here:
//For simplicity,assume MAX_SIZE is exactly divisible by the number of
threads.
int num_threads = atoi(argv[1]);

struct thread_param
{
    int start; int end;
};

int chunk = MAX_SIZE/num_threads;

std::vector<std::thread> workers(num_threads);
std::atomic<int> histogram[RANGE];

for(int i=0;i<RANGE;i++) histogram[i].store(0,memory_order_seq_cst);

for(int i=0;i<num_threads;i++)
{
    thread_param tp;
    tp.start = i*chunk;
    tp.end = std::min((i+1)*chunk,MAX_SIZE);
    std::thread tmp{calc_histogram,tp}
    workers[i] = std::move(tmp);
}
```



```
for(int i=0;i<num_threads;i++) workers[i].join();
```

```
void calc_histogram(thread_param tp)
```

```
{  
    for(int i=tp.start;i<tp.end;i++)  
    {  
        int pos = s[i];  
        histogram[pos].fetch_and_add(1,memory_order_seq_cst);  
    }  
}
```

5. MPI_Thread types

(10 pts)

Describe the difference between MPI_Thread_Single, MPI_Thread_Funneled and MPI_Thread_Multiple. Provide a use case for MPI_Thread_Multiple (the pseudocode or description is sufficient). What is the advantage of using MPI_Thread_Multiple in your use case?

MPI_Thread_Single : One thread executes

MPI_Thread_funneled : Main thread makes MPI calls

MPI_Thread_multiple : multiple threads can make MPI calls

The halo exchange example discussed in lecture is a possible use case.

6. MPI - Implementation of Broadcast

(10+5 pts)

Implement a broadcast-like operation on a group of processes. Assume the operation is initiated on the process with rank 0, and the data needs **to reach all other processes**. The broadcast should use a tree algorithm, where the processes are arranged in a binary tree, each process sends data to its children. Children of a process with rank i are the processes with rank $2 * i + 1$ and $2 * i + 2$. E.g., process 0 first sends data to 1 and 2, then 1 sends it to 3 and 4, while 2 sends to 5 and 6, and so on.

```
const int ASIZE = ...;
int size, rank;
MPI_Init(NULL, NULL);
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);

int bArray[ASIZE];
if(rank == 0) {
    for (int i = 0; i < ASIZE; ++i) {
        bArray[i] = computeithElement(i);
    }
}

// Implement your broadcast algorithm here to send bArray from process
// 0 to all the other processes in MPI_COMM_WORLD.

int nbr1 = 2*i+1;
int nbr2 = 2*i+2;
int p = (rank-1)/2;

int nbr_count = 0;
if(nbr1 < size)
{
    MPI_Send(&bArray[0], ASIZE, MPI_INTEGER, nbr1, 123, MPI_COMM_WORLD);
}
if(nbr2 < size)
{
    MPI_Send(&bArray[0], ASIZE, MPI_INTEGER, nbr2, 123, MPI_COMM_WORLD);
}
```

```

if(rank > 0)
{
    MPI_Request r;
    MPI_Status s;
    MPI_Irecv(&bArray[0], ASIZE, MPI_INTEGER, p, 123, MPI_COMM_WORLD, r);
    MPI_Wait(r, s);
}

```

```

MPI_Finalize();

```

B. What is the communication cost of your implementation? You can assume a communication model where the cost of send/rcv is $\alpha + \beta * m$ where α is the latency, β is the cost per byte and m is the message size in bytes.

Number of stages = $\log P$

Communication cost = $\log P * (\alpha + 4 * \beta * ASIZE)$

