

# PDES


(Don't worry, this isn't calculus)

# **Parallel Discrete Event Simulation**

# Classification of Computer Simulations

**Computer Simulation**


# Classification of Computer Simulations



# Classification of Computer Simulations

Computer Simulation		
Continuous	Discrete	
	Time stepped	Event driven

# Classification of Computer Simulations



# Discrete Event Simulations

- Logical Processes (LPs) execute events
- Executing an event updates the LP's state
- Events have a virtual timestamp
- Events must be executed in order

# Applications

- Traffic analysis
- Military battles
- Networks
- Circuits
- Economic models
- and many more...






# Implementation

- Single event queue
- Sorted by timestamp
- Loop over queue and execute events
- Efficiency depends on queue used
- Very simple

# Discrete Event Simulations

LPs






Event List	
Timestamp	Destination
6	A
12	B
13	B
17	C

Sim Time
6

# Discrete Event Simulations

LPs






Event List	
Timestamp	Destination
6	A
12	B
13	B
17	C

Sim Time
6

# Discrete Event Simulations

LPs



Event List	
Timestamp	Destination
12	B
13	B
15	C
17	C

Sim Time
12

# Pseudocode

```
while (running) {  
    Event* e = eventList.pop();  
    LP* lp = e->destination();  
    lp->execute(e);  
}
```

Pretty simple right?

# Limitations of Sequential DES

- Millions of LPs
- Billions of events to simulate
- Sequential simulations will take too long

# How will parallelization help?

- Events are generally very small
- Simulations can have of billions of events

**Benefits of parallelism come from executing many events at once**

# How do we parallelize DES?

- Distribute LPs across processors
- Each processor has its own event list and virtual clock
- How to synchronize event lists and clocks?



# (Super) Naive Implementation

```
while (running) {  
    AllReduce(simTime);  
    if (events.top()->ts==simTime) {  
        Event* e = events.pop();  
        LP* lp = e->destination();  
        lp->execute(e);  
    }  
}
```


# **(Super) Naive Implementation**

- Exactly matches sequential semantics
- NO PARALLELISM
- Way too much synchronization

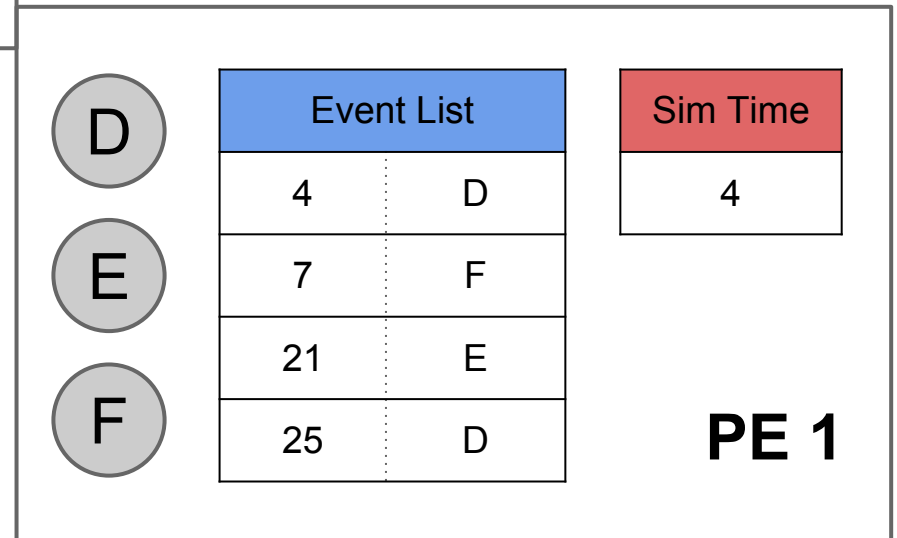
**Moral of the Story:**

**We need to relax our ordering restrictions!**


# What is the fundamental problem?



**Can PE 0 safely execute any events?**




# What is the fundamental problem?

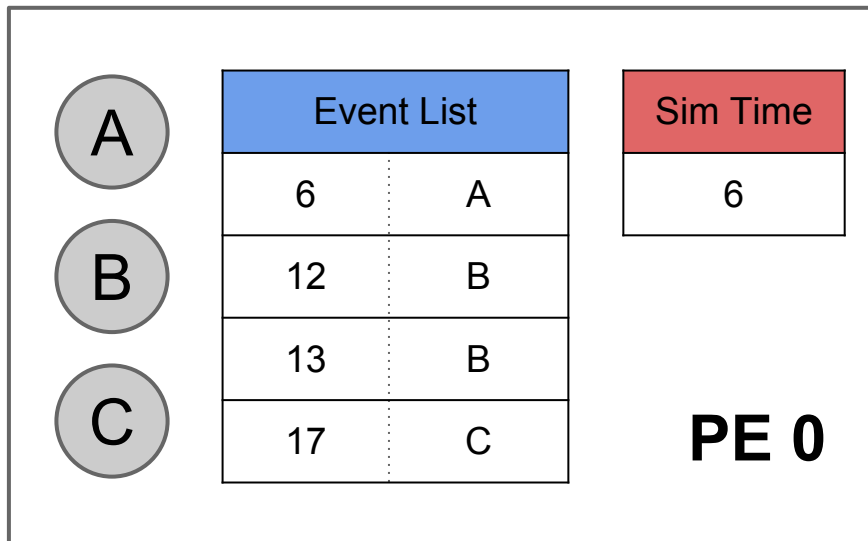


**Can PE 0 safely execute any events?**

**What if the first event on PE 1 generates event (5,A)?**



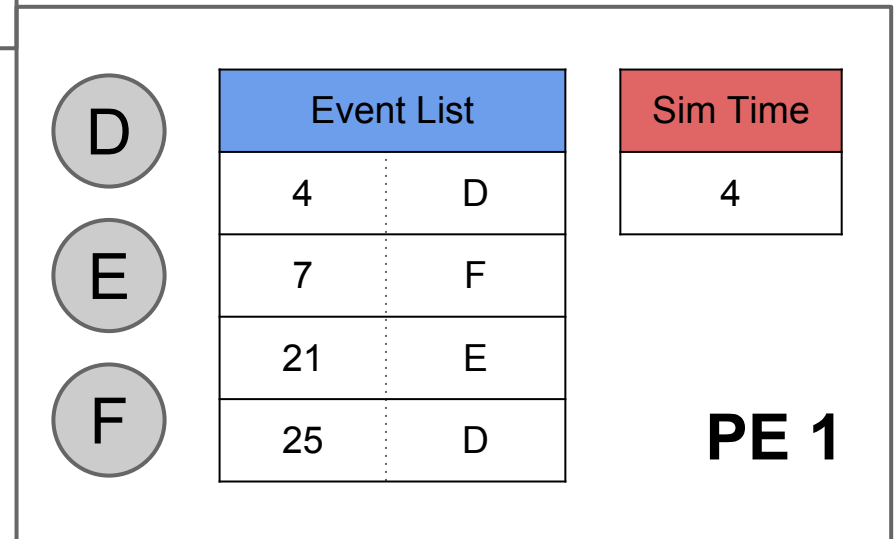
# What is the fundamental problem?



**Can PE 0 safely execute any events?**

**What if the first event on PE 1 generates event (5,A)?**

**Causality Error!**



What is the fundamental problem?

**How can we  
ever execute  
events!?**

What if the first event  
on PE 1 generates  
event (5,A)?

Can PE 0 safely  
execute any  
events?

A

B

C

Event List	
6	
12	B
13	B
14	C

Sim Time
6

D

E

F

Event List	
4	D
7	F
21	E
25	D

Sim Time
4

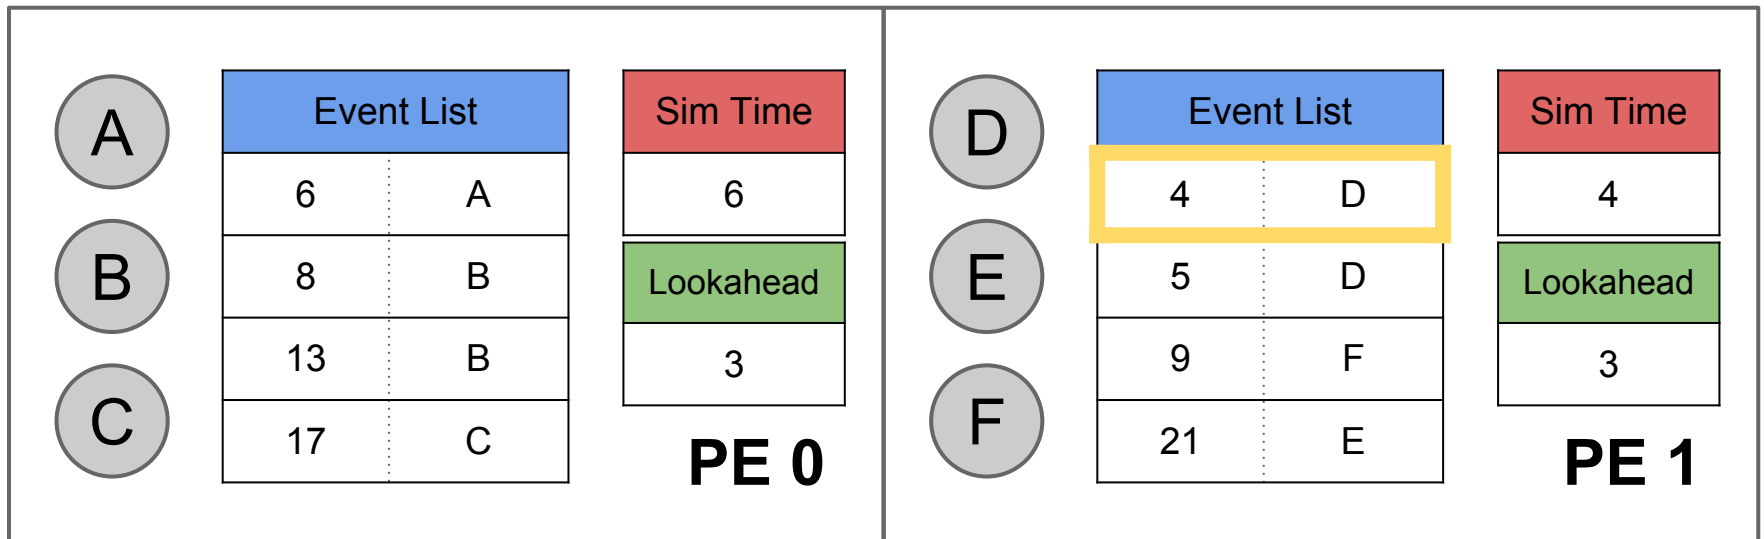
PE 1

# Two Approaches

- Conservative (don't screw up)
  - Only execute events we know won't be preempted
  - High amount of synchronization
  - Low amount of parallelism
  - Inflexible
  - Simple
- Optimistic (if we screw up, we'll fix it)
  - Execute events freely
  - If a causality error occurs, rollback the processor
  - High amount of parallelism
  - Flexible
  - Complex

# Conservative (Windowed)

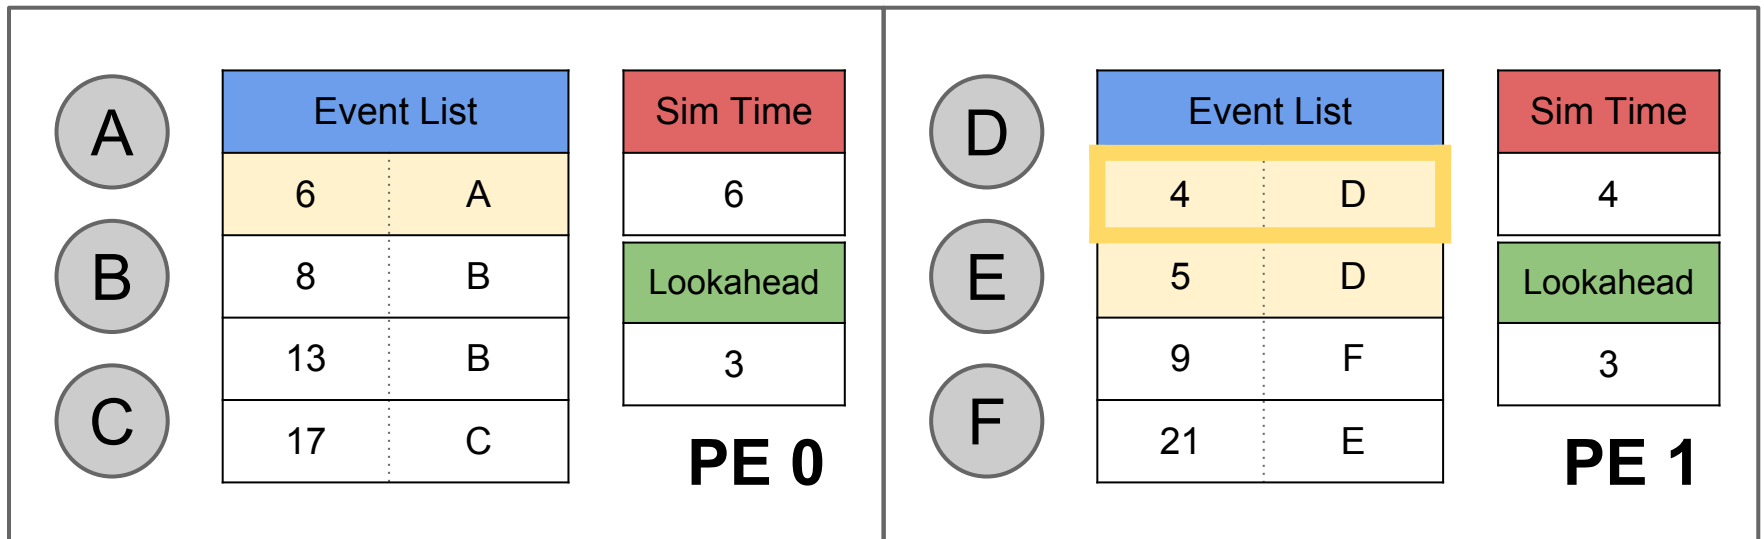
- Requires a model-specific lookahead
- Determine the min timestamp in the system
- Execute events in (min + lookahead) window





# Conservative (Windowed)

- Requires a model-specific lookahead
- Determine the min timestamp in the system
- Execute events in (min + lookahead) window



# (Super) Naive Implementation

```
while (running) {  
    AllReduce(simTime);  
    if (events.top()->ts==simTime) {  
        Event* e = events.pop();  
        LP* lp = e->destination();  
        lp->execute(e);  
    }  
}
```

# Windowed Implementation

```
while (running) {  
    AllReduce(simTime);  
    while (events.top()->ts <  
           simTime + lookahead) {  
        Event* e = events.pop();  
        LP* lp = e->destination();  
        lp->execute(e);  
    }  
}
```

# Windowed Analysis

- Very similar to naive solution
- Performance depends on lookahead
- Low lookahead = low parallelism and high synchronization
- Workload can be unbalanced

# Optimistic

- Execute events freely
- When an event is received with a smaller timestamp than your clock, rollback
- How do we rollback efficiently?
- How do rollbacks affect performance?
- How many PEs will a rollback affect?

# Rollbacks

- Save previous events (How many?)
- Revert your own state (How?)
- Cancel sent events (How?)

# Saving Previous Events

- Events take up memory
  - Limits how many events we can save
  - Need to reclaim memory periodically
- What can we safely reclaim?
- Find the Global Virtual Time (GVT)
  - Minimum clock time of the system
  - Everything prior to this can be reclaimed
  - Events with observable effects can be committed

# GVT

- Global synchronization required
- All events must be accounted for
- Can be synchronous or asynchronous



# Reverting Your State

- State saving
  - Save the states of LPs after each event
  - Rolling back is equivalent to reverting states
  - High memory consumption
  - Need to reclaim memory more often
- Reverse computation
  - During rollback execute events in reverse
  - Better for memory
  - Overhead of executing in reverse
  - Reverse computation is complex
  - Can compilers help?

# Cancelling Events

- Events sent erroneously must be cancelled
- First we must find the event
  - If it's local, that's easy
  - If it's remote we need to send an anti-event
- Then we must cancel it
  - If they weren't executed, just delete them
  - If they have been executed, do a rollback
- Rollbacks can snowball out of control

# Pseudocode

```
while (running) {  
    while (executing_events) {  
        check_for_rollbacks();  
        Event* e = events.pop();  
        LP* lp = e->destination();  
        lp->execute(e);  
    }  
    compute_gvt();  
}
```

# Summary

## Two Main Classes of PDES:

- Conservative

- Low parallelism/High synchronization cost
- Model dependent
- Low memory footprint

- Optimistic

- High parallelism/Low synchronization cost
- Model independent
- Memory Hungry
- Rollbacks can snowball